

# Distance metrics

# Continuous variables

## Minkowski family of distances

$$D(i, j) = {}^{1/p}\sqrt{\sum_{k=1}^N |x_{ik} - x_{jk}|^p}$$

N features (dimensions)

$$D(i, j) > 0$$

$$D(i, i) = 0$$

properties

$$D(i, j) = D(j, i)$$

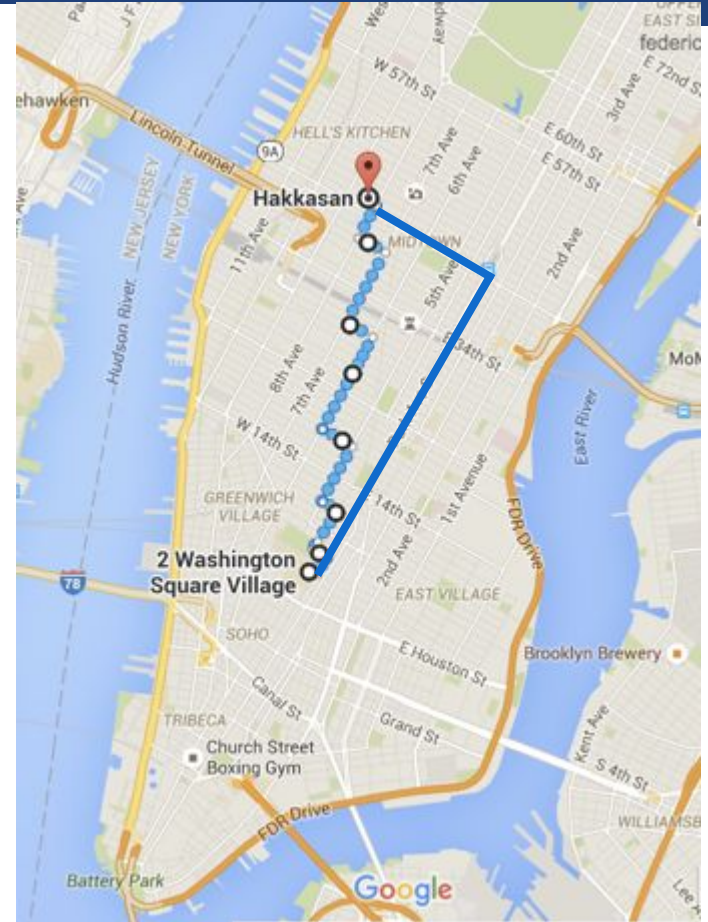
$$D(i, j) \leq D(i, k) + D(k, j)$$

# Continuous variables

## Minkowski family of distances

Manhattan:  $p=1$

$$D_{Man}(i, j) = \sum_{k=1}^N |x_{ik} - x_{jk}|$$

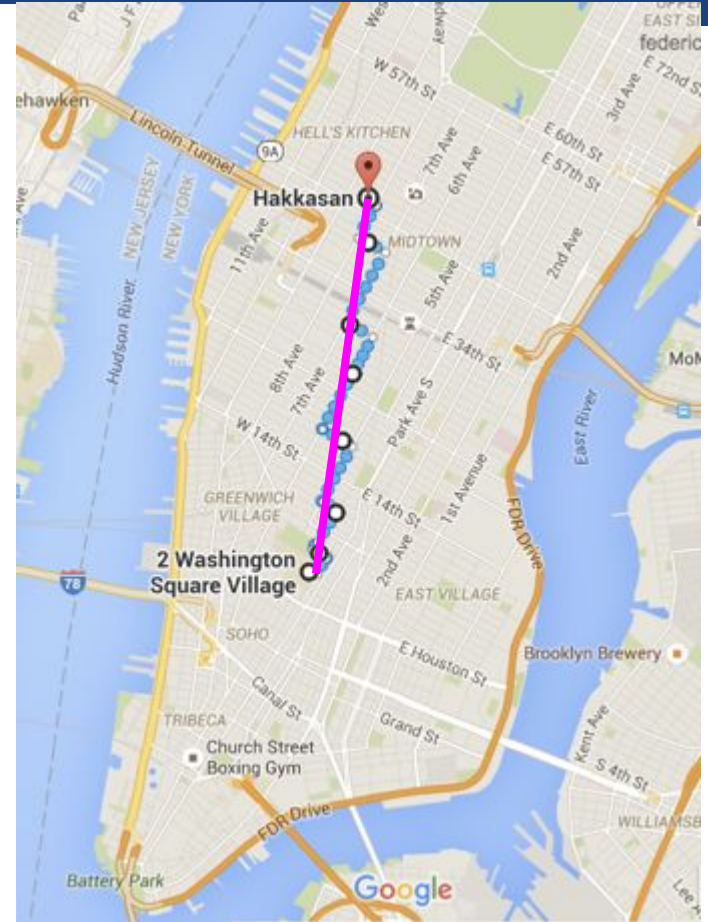


# Continuous variables

## Minkowski family of distances

Euclidean:  $p=2$

$$D_{Euc}(i, j) = \sqrt{\sum_{k=1}^N |x_{ik} - x_{jk}|^2}$$



# Continuous variables

## Great Circle distance

$$D(i, j) = R \arccos (\sin \phi_i \cdot \sin \phi_j + \cos \phi_i \cdot \cos \phi_j \cdot \cos \Delta \lambda)$$

features

latitude and longitude

$$\phi_i, \lambda_i, \phi_j, \lambda_j$$



# Categorical variables: binary

Uses presence/absence of features in data

$M_{i=0,j=0}$  : number of features in neither  
 $M_{i=1,j=1}$  : number of features in both  
 $M_{i=1,j=0}$  : number of features in  $i$  but not  $j$   
 $M_{i=0,j=1}$  : number of features in  $j$  but not  $i$

		observation $i$		
		1	0	sum
observation $j$	1	M11	M10	M11+M10
	0	M01	M00	M01+M00
	sum	M11+M01	M10+M00	M11+M00+ M01+ M10



# Categorical variables: binary

Uses presence/absence of features in data

## Simple Matching Coefficient or Rand similarity

$$SMC(i, j) = \frac{M_{i=0,j=0} + M_{i=1,j=1}}{M_{i=0,j=0} + M_{i=1,j=0} + M_{i=0,j=1} + M_{i=1,j=1}}$$

		observation <i>i</i>		
		1	0	sum
observation <i>j</i>	1	M11	M10	M11+M10
	0	M01	M00	M01+M00
	sum	M11+M01	M10+M00	M11+M00+M01+ M10

## Simple Matching Distance

$$SMD(i, j) = 1 - SMC(i, j)$$

$M_{i=0,j=0}$  : number of features in neither  
 $M_{i=1,j=1}$  : number of features in both  
 $M_{i=1,j=0}$  : number of features in *i* but not *j*  
 $M_{i=0,j=1}$  : number of features in *j* but not *i*

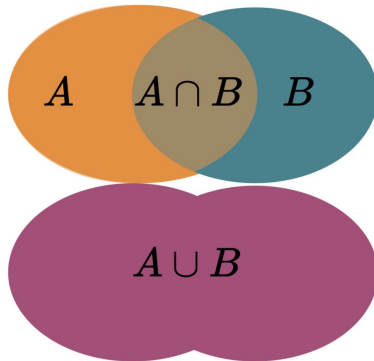
# Categorical variables: binary

## Jaccard similarity

$$J(i, j) = \frac{M_{i=1, j=1}}{M_{i=0, j=1} + M_{i=1, j=0} + M_{i=0, j=0}}$$

## Jaccard distance

$$D(i, j) = 1 - J(i, j)$$



		observation <i>i</i>		
		1	0	sum
observation <i>j</i>	1	M11	M10	M11+M10
	0	M01	M00	M01+M00
	sum	M11+M01	M10+M00	M11+M00+M01+M10

$M_{i=0, j=0}$  : number of features in neither  
 $M_{i=1, j=1}$  : number of features in both  
 $M_{i=1, j=0}$  : number of features in *i* but not *j*  
 $M_{i=0, j=1}$  : number of features in *j* but not *i*



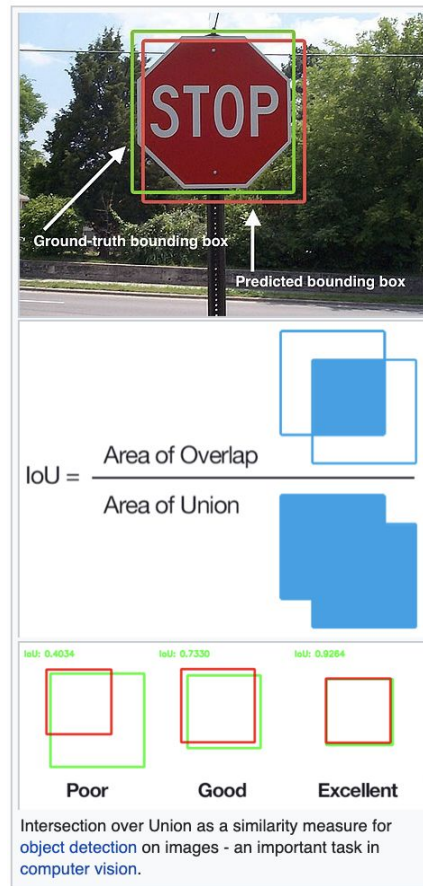
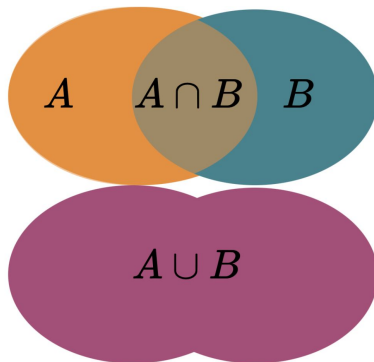
# Categorical variables: binary

## Jaccard index

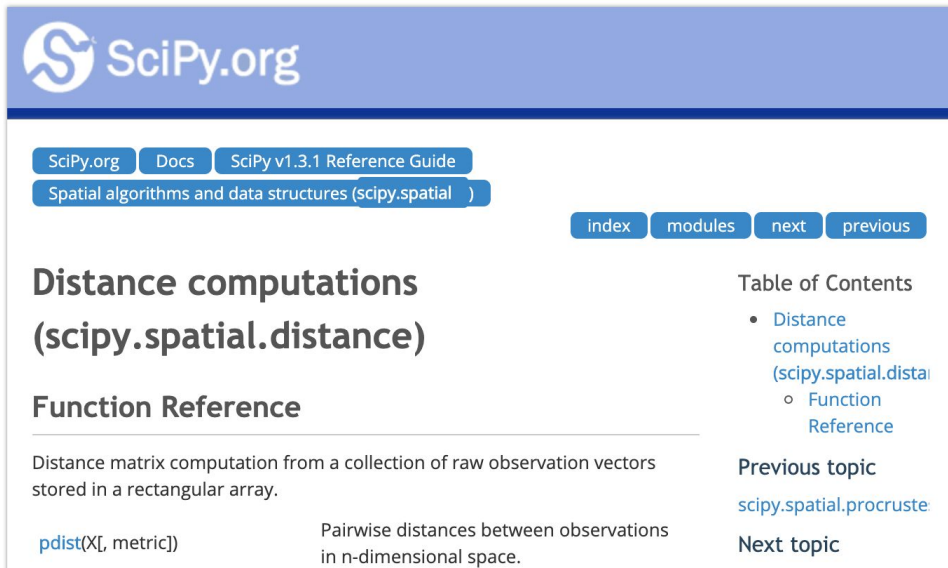
$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Application to Deep Learning  
for image recognition

Convolutional Neural Nets



# Another useful package for scientific Python: SciPy



The screenshot shows the SciPy.org website. At the top is the SciPy.org logo. Below it are navigation links: SciPy.org, Docs, and SciPy v1.3.1 Reference Guide. A sub-link for 'Spatial algorithms and data structures (scipy.spatial)' is highlighted. On the right, there are buttons for 'index', 'modules', 'next', and 'previous'. The main content area is titled 'Distance computations (scipy.spatial.distance)' and 'Function Reference'. It describes 'Distance matrix computation from a collection of raw observation vectors stored in a rectangular array.' and lists the function `pdist(X[, metric])` with the description 'Pairwise distances between observations in n-dimensional space.' On the right side of the page, there is a 'Table of Contents' with a link to 'Distance computations (scipy.spatial.distance)' and a sub-link to 'Function Reference'. Below that are links for 'Previous topic' (scipy.spatial.procruste) and 'Next topic'.

<https://docs.scipy.org/doc/scipy/reference/spatial.distance.html>

Distance functions between two boolean vectors (representing sets) `u` and `v`. As in the case of numerical vectors, `pdist` is more efficient for computing the distances between all pairs.

`dice(u, v[, w])`

Compute the Dice dissimilarity between two boolean 1-D arrays.

`hamming(u, v[, w])`

Compute the Hamming distance between two 1-D arrays.

`jaccard(u, v[, w])`

Compute the Jaccard-Needham dissimilarity between two boolean 1-D arrays.

Distance functions between two numeric vectors `u` and `v`. Computing distances over a large collection of vectors is inefficient for these functions. Use `pdist` for this purpose.

`braycurtis(u, v[, w])`

Compute the Bray-Curtis distance between two 1-D arrays.

`canberra(u, v[, w])`

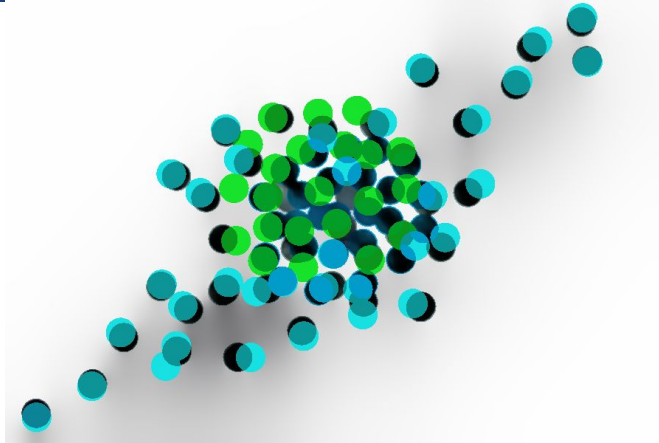
Compute the Canberra distance between two 1-D arrays.

`chebyshev(u, v[, w])`

Compute the Chebyshev distance.

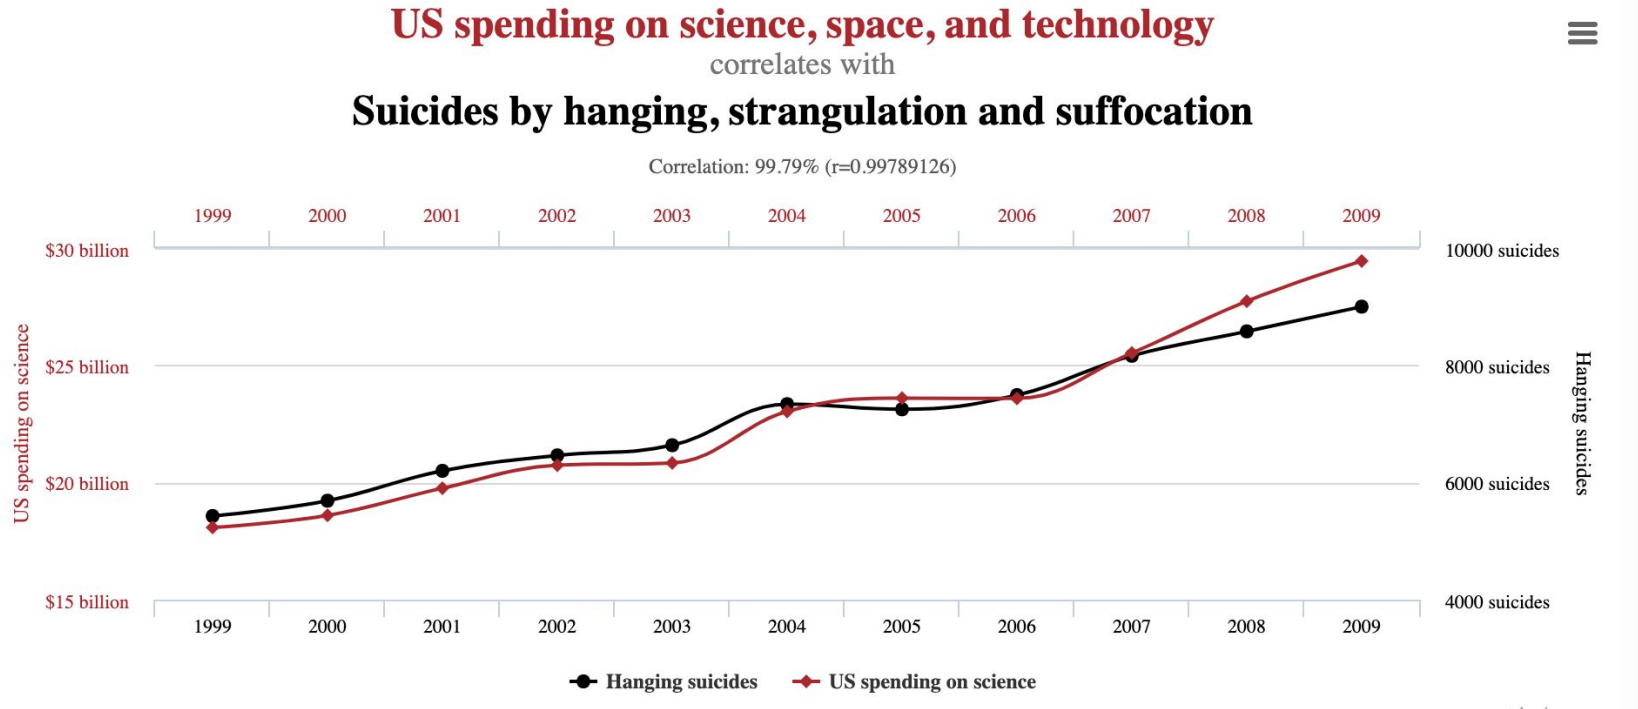
`cityblock(u, v[, w])`

Compute the City Block (Manhattan) distance.



# Data whitening

# Data can have covariance (and it almost always does!)

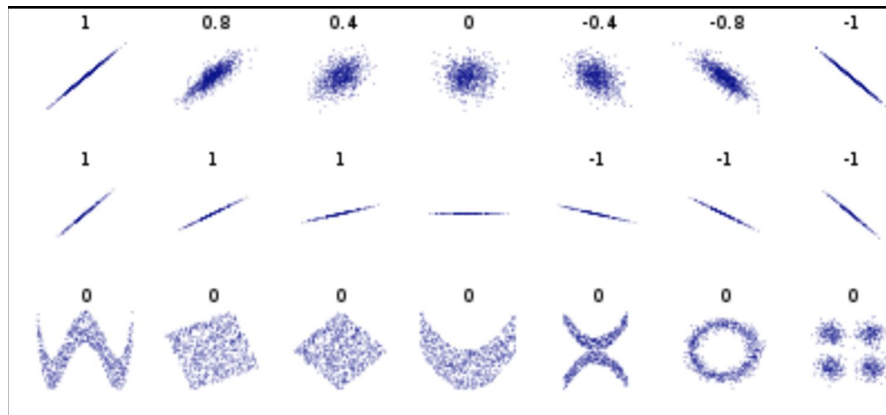


<https://www.tylervigen.com/spurious-correlations>

# Data can have covariance (and it almost always does!)

Pearson's correlation (linear correlation)

$$r_{xy} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$



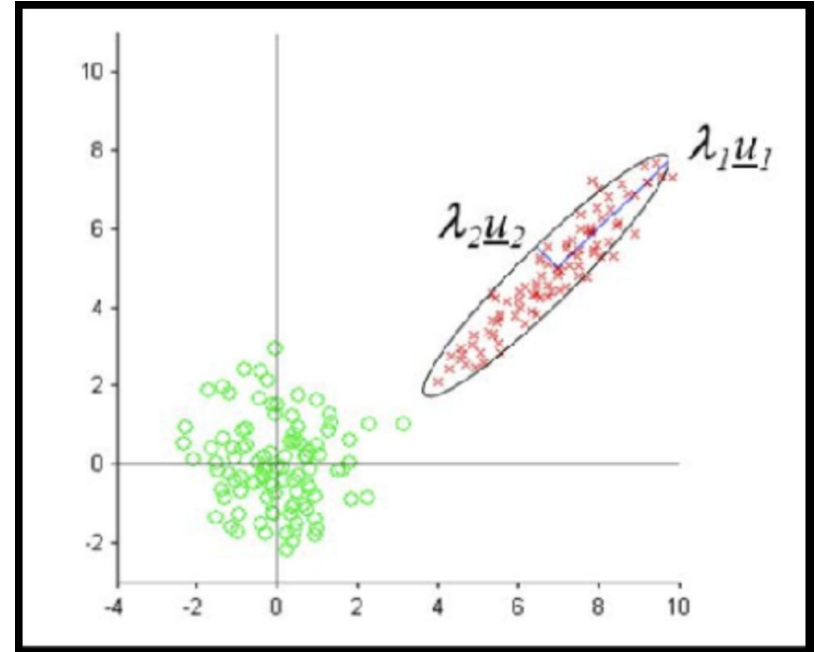
Pearson's correlation = covariance/ product of the two standard deviations

# Generic preprocessing

Original data

Data that is not correlated appear as a sphere in the N-dimensional feature space

Standardized data



# Generic preprocessing

for each feature: divide by  
standard deviation and subtract  
mean

Original data

mean of each feature should be  
0, standard deviation of each  
feature should be 1

Standardized data

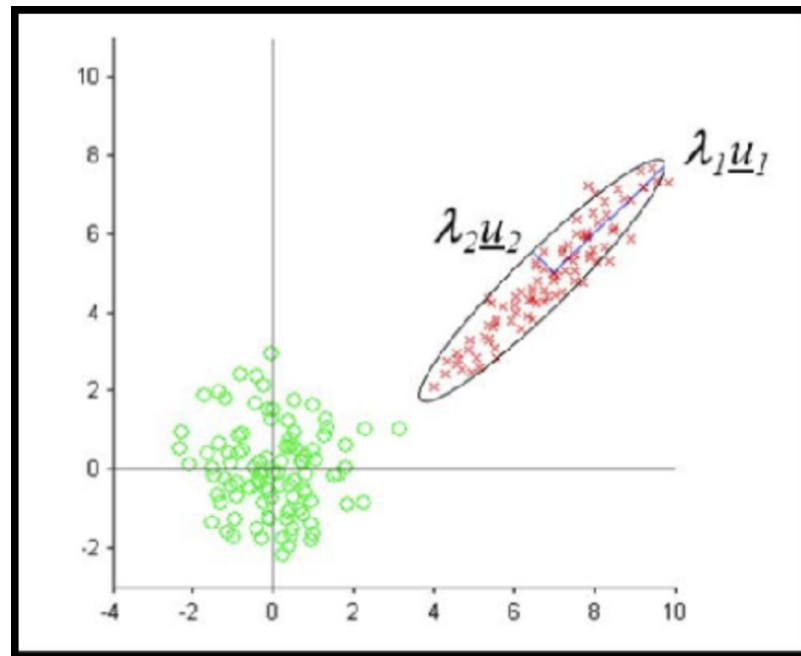
```
X = preprocessing.scale(X, axis=0)
Last executed 2018-12-12 09:35:39 in 46ms

X.mean(axis=0)
Last executed 2018-12-12 09:35:40 in 13ms

array([[ 3.85590369e-16, -6.93196168e-17, -5.90549813e-16, -5.95882091e-16,
        -8.49165306e-16, -1.57568821e-15, -8.00508267e-16,  5.55890004e-16,
        -5.16564452e-16,  1.09378357e-15,  3.46598084e-16,  2.31954102e-16,
         2.78611537e-16, -2.51283611e-16,  8.66495210e-18,  3.03939858e-16,
        -3.66594127e-17, -9.27149875e-16, -6.39873386e-16,  2.93275302e-17,
         9.19817992e-17,  6.33208038e-18, -1.99960433e-17,  9.55144336e-16,
        -2.20623011e-16,  6.93196168e-17, -9.46479383e-17,  2.26621824e-16])

X.std(axis=0)
Last executed 2018-12-12 09:36:28 in 19ms

array([[1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.]])
```





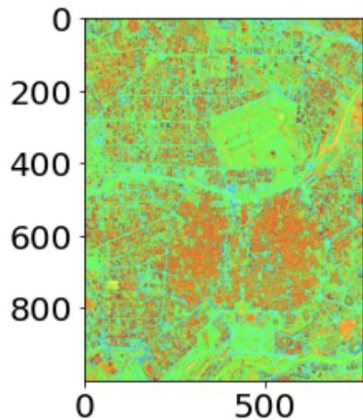
# Generic preprocessing

For image processing (e.g. segmentation) often you need to min-max preprocess

```
from sklearn import preprocessing
Xopscaled = preprocessing.minmax_scale(op.reshape(op.shape[1] * op.shape[0], 3).astype(float), axis=1)
Xopscaled.reshape(op.shape) [200, 700]
```

```
pl.imshow(Xopscaled.reshape(op.shape));
```

Clipping input data to the valid range for imshow with RGB data





# Clustering

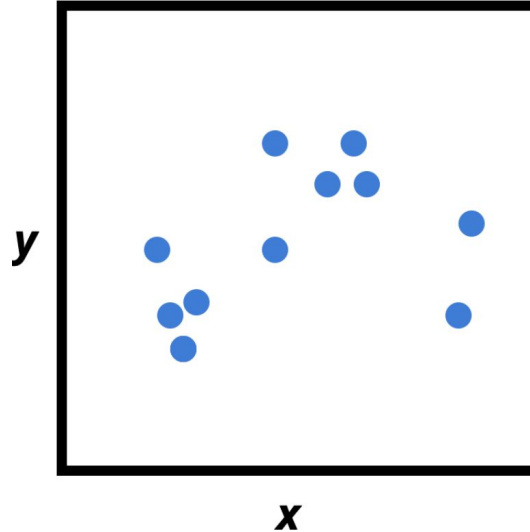
# Clustering is an unsupervised learning method

GOAL: partitioning data in maximally homogeneous,  
maximally distinguished subsets.

observed

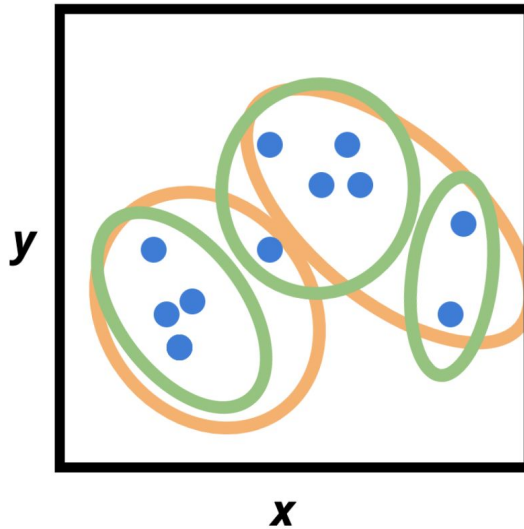
**features:**

$(x, y)$



# Clustering is an unsupervised learning method

all features are  
observed for all  
objects in the  
sample  
 $(x, y)$



how should I  
group the  
observations in  
this feature  
space?

*e.g.: how  
many groups  
should I  
make?*

# Optimal clustering



<http://www-bcf.usc.edu/~soltanol/Applications.html>

## internal criterion:

members of the cluster should be similar to each other (intra cluster compactness)

## external criterion:

objects outside the cluster should be dissimilar from the objects inside the cluster



**Tigers**



**Whales**



**raptors**

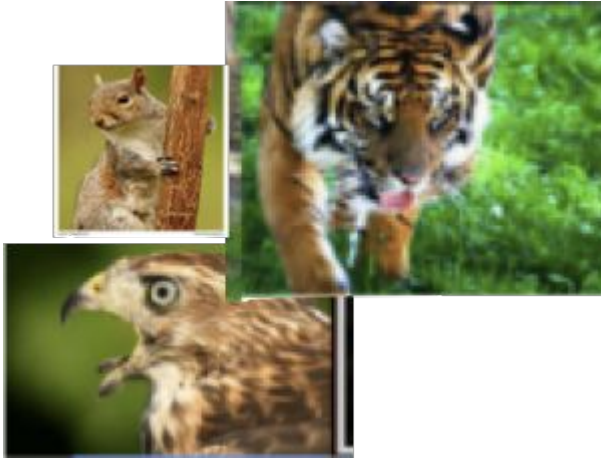
**zoologist's clusters**

## internal criterion:

members of the cluster should be similar to each other (intra cluster compactness)

## external criterion:

objects outside the cluster should be dissimilar from the objects inside the cluster



yellow/green

photographer's clusters



black/white/blue



# Optimal clustering

**the optimal clustering depends on:**

- how you define similarity/distance
- the purpose of the clustering

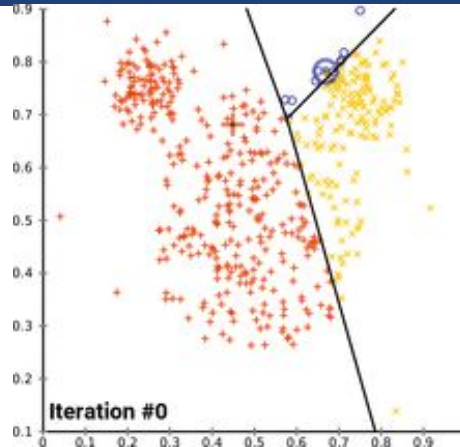
# Optimal clustering

**The ideal clustering algorithm should have:**

- Scalability
- Ability to deal with different types of attributes
- Discovery of clusters with arbitrary shapes
- Minimal requirement for domain knowledge
- Deals with noise and outliers
- Insensitive to order
- Allows incorporation of constraints
- Interpretable

# How does clustering work?

- **Partitional**
  - **Hard clustering:** K-means (McQueen '67), K-medoids (Kaufman & Rausseeuw '87)
  - **Soft Clustering:** Expectation Maximization (Dempster, Laird, Rubin '77)
- **Hierarchical**
  - **Agglomerative (bottom-up)**
  - **Divisive (top-down)**
- **also:**
  - **Density based**
  - **Grid based**
  - **Model based**



# Clustering by partitioning

# K-means

Choose N “centers” guesses: random points in the feature space

**repeat:**

- Calculate distance between each point and each center
- Assign each point to the closest center
- Calculate the new cluster centers

**until (convergence):**

when clusters no longer change

# K-means

**Objective:** minimizing the aggregate distance within the cluster.

**Order:** #clusters #dimensions #iterations #datapoints  **$O(KdN)$**

## CONs:

- **Its non-deterministic:** the result depends on the (random) starting point
- **It only works where the mean is defined:** alternative is K-medoids which represents the cluster by its central member (median), rather than by the mean
- **Must declare the number of clusters upfront** (how would you know it?)

## PROs:

**Scales linearly with d and N**

# K-means

**Objective:** minimizing the aggregate distance within the cluster.

**Order:** #clusters #dimensions #iterations #datapoints  **$O(KdN)$**

**$O(KdN):$**

complexity scales linearly with:

- $d$  number of dimensions
- $N$  number of datapoints
- $K$  number of clusters



# K-means: the objective function

**Objective:** minimizing the aggregate distance within the cluster.

$$\text{total intra-cluster variance} \quad \sum_k \sum_{i \in k} (\vec{x}_i - \vec{\mu}_k)^2$$

**Order:** #clusters #dimensions #iterations #datapoints  $O(KdN)$

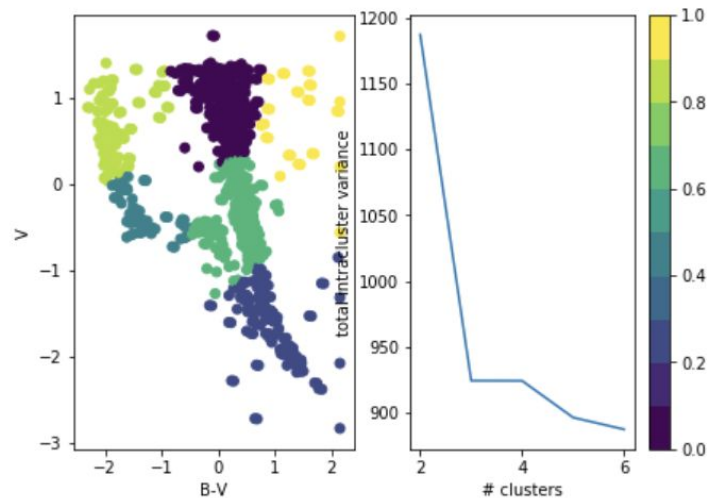
## Must declare the number of clusters

either you know it because of domain knowledge

or

you choose it after the fact: "*elbow method*"

```
for i in range(1, nmaxc):
    c = KMeans(n_clusters=i, random_state=302).fit(X2)
    nc.append(c.inertia_)
    print("i.c. variance with {} clusters {:.2f}".format(i, c.inertia_))
plt.plot(range(1, nmaxc), nc)
```



# K-means: the objective function

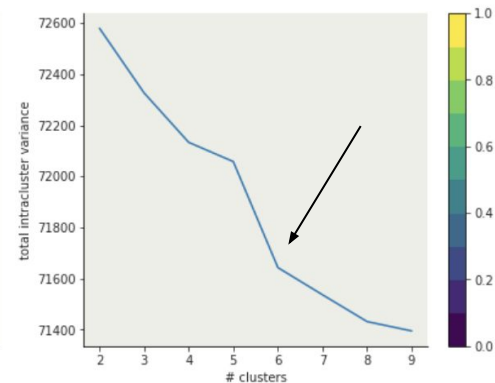
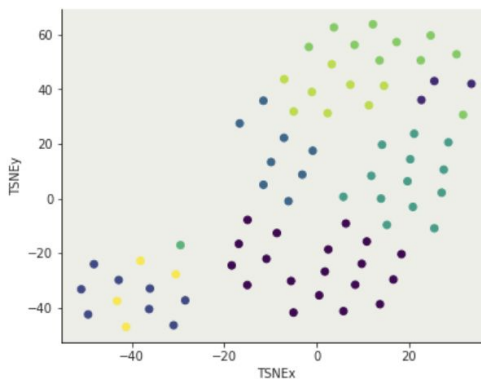
**Objective:** minimizing the aggregate distance within the cluster.

total intra-cluster variance 
$$\sum_k \sum_{i \in k} (\vec{x}_i - \vec{\mu}_k)^2$$

**Order:** #clusters #dimensions #iterations #datapoints **O(KdN)**

**Must declare the number of clusters upfront** (how would you know it?)

either domain knowledge or  
after the fact: "elbow method"



# K-means hyperparameters

```
class sklearn.cluster. KMeans (n_clusters=8, init='k-means++', n_init=10, max_iter=300, tol=0.0001,  
precompute_distances='auto', verbose=0, random_state=None, copy_x=True, n_jobs=None, algorithm='auto') ¶
```

[\[source\]](#)

- *n\_clusters* : number of clusters
- *init* : the initial centers or a scheme to choose the center
  - 'k-means++' : selects initial cluster centers for k-mean clustering in a smart way to speed up convergence.
  - 'random': choose k observations (rows) at random from data for the initial centroids.
  - If an ndarray is passed, it should be of shape (n\_clusters, n\_features) and gives the initial centers.
- *n\_init* : if >1 it is effectively an ensemble method: runs n times with different initializations
- *random\_state* : for reproducibility

# Convergence criteria

## General

Any time you have an objective function (or loss function) you need to set up a tolerance : if your objective function did not change by more than  $\varepsilon$  since the last step you have reached convergence (i.e. you are satisfied)

$\varepsilon$  is your tolerance

## For clustering:

convergence can be reached if  
*no more than  $n$  data point changed cluster*

$n$  is your tolerance